

Differentially Private Interactive Media

JAN VAN BRUGGEN
jvanbrug@caltech.edu

CUTTER CORYELL
ccoryell@caltech.edu
California Institute of Technology

JAMES CHANG
jcchang@caltech.edu

1. Introduction

In the modern era, entities who own and sell digital content are increasingly concerned with unauthorized access to that content. A variety of “digital rights management” (DRM) technologies have been introduced to attempt to restrict content access to authorized users, with mixed results. In many cases arms races have resulted between content-providers and users in terms of rights management technologies and methods for breaking them. These technologies usually make it hard to share content between authorized users and unauthorized users and are commonly criticized for inconveniencing authorized users — some content-providers push their lack of DRM as a selling point, and a general anti-DRM movement has sprung up.

One way to prevent unauthorized sharing of content is to change the nature of the content to make it fundamentally unsharable. For example, by *streaming* content as opposed to handing out content in bundles which can be saved locally, it becomes non-trivial for the regular user to share their received content. Many music, movie/television, and even video-game services take this approach to control content access. However, it is not difficult for an unauthorized user to record the data stream in a way that can be saved locally, such that he manufactures his own content bundles which he can share. We say that the user has *fully reconstructed* the content if she can provide an identical experience to an unauthorized user as an authorized user would receive through the official content channel.

It becomes significantly harder for a user to fully reconstruct the original content if the streamed content is in some way interactive. For example, in a streamed video-game, the data that is streamed to the user depends on the the input that the user provides. In order to provide the same experience to unauthorized users as the one that authorized users receive, the reconstructor must provide every possible time series of inputs to the content provider and record every resulting stream. If the content is deterministic — the same time series of inputs always leads to the same data stream — then this reconstruction ensures that any inputs that authorized users can give can also be given by unauthorized users *with the same results*. This satisfies our definition of full reconstruction, because the experience of the unauthorized user

is identical to that of the authorized user.¹

Yet another change that can be made to how content is provided to make reconstruction more difficult is the addition of noise. This method is motivated by differential privacy, which is a notion of how the probability of a certain output to a query will change given a change in the queried database. By adding noise to a query in a special way, the particular existence or absence of a datum in the database can be concealed while still revealing useful information that respects differential privacy. The general idea of using noise to make it harder to reconstruct content is one we would like to borrow. With noise,² every stream a content provider serves to a user is a random variable. Someone attempting to reconstruct the original content who, as before, only samples from the stream once (or once for every possible input time series if the content is interactive) will provide users with an empirical probability distribution where only one output has nonzero probability: that particular output the reconstructor sampled from the original content provider. This is quite possibly a poor approximation to the actual probability distribution governing possible outputs in the official stream.

How interactivity and noise make reconstruction harder, both individually and together, is an interesting question that we attempt to answer.

2. Methods

To investigate how the ease of reconstruction of a specific content-base depends on the interactiveness of the content stream and on the noise added to the content stream, we set up two simple models. In the first model, which lacks interaction, there is a space of possible outputs with utilities assigned by the content-creator — the particular output given to a user is chosen by the *exponential mechanism* from the differential privacy literature.

The output for the second model is a Markovian sequence of items. The probability of any particular item appear at a certain position in the sequence is determined entirely from the previous item in the

¹We define two users to have *identical experience* if, given the same time series of user inputs, the true probability distributions of certain content being streamed to each user is the same. This will be further motivated when we discuss noise in streaming.

²To be clear, when we refer to *noise* we mean simply the introduction of randomness in the potential stream output — far more general than the literal addition of noise values to the data.

sequence. The first item of the sequence is chosen uniformly randomly, and subsequent items are chosen via the exponential mechanism. As in the first model, the content-creator assigns utilities to each item-to-item transition which are used by the exponential mechanism. However, in this model we allow for interaction. At query-time, the user provides a vector of preferences, one for each item in the item space. The utilities of the transitions away from a particular item are allowed to be functions of the user preference value for that item.

We now develop the two models mathematically.

2.1. Model 1

Let \mathcal{S} be the set of all possible outputs to a streaming query. Let $\mathcal{I} \subseteq \mathcal{S}$ be a set of *intended* outputs — these being, informally, the outputs that the content-creator actually wishes the users to receive in the absence of noise. The content creator specifies $\mathbf{u} = [(u_1, s_1), (u_2, s_2), (u_3, s_3), \dots]$, where each $u_j \in (0, 1]$ is a utility and each $s_j \in \mathcal{S}$. Here, \mathbf{u} is a multiset (elements can occur more than once) with any desired cardinality. In terms of differential privacy, \mathbf{u} is a database. We have $\mathcal{I} = \mathcal{I}(\mathbf{u}) \equiv \{s \in \mathcal{S} \mid \exists u \text{ s.t. } (u, s) \in \mathbf{u}\}$.

The sampling mechanism is defined as follows:

$sample_1(\mathbf{u}, \varepsilon) :$

Output $s \in \mathcal{S}$ with probability $\propto \exp(\frac{1}{2}\varepsilon U(s))$

where $U(s) \equiv u_j$ if $s = s_j \in \mathcal{I}$ or $U(s) = 0$ if $s \notin \mathcal{I}$.

This mechanism is the exponential mechanism of differential privacy with privacy parameter ε where the utility function is defined over the outcome space of all possible outputs, including unintended outputs. The sensitivity of $U(s)$ is 1 because $U(s) \in [0, 1]$ and is implicitly in the mechanism definition. Because it is just an application of the exponential mechanism, $sample_1$ is ε -differentially private. What this means intuitively is that the change in probability of a certain output being sampled if the output (or any other output, in fact) changes from intended to unintended (or back) cannot be more than a certain factor — e^ε to be precise.

The interesting question to be answered by this model is how the reconstruction quality of \mathbf{u} changes with ε and with the number of queries to $sample_1(\mathbf{u}, \varepsilon)$. What the reconstructor, which we shall refer to as the adversary henceforth, tries to build is a probability distribution over \mathcal{S} , $p_{\text{approx}}(s)$ which approximates the true probability distribution $p(s)$ over outputs of $sample_1(\mathbf{u}, \varepsilon)$. There are many conceivable ways to generate $p_{\text{approx}}(s)$ from a finite number of samples — the most straightforward of which is constructing an empirical distribution where the probability of s is equal to the proportion of all outputs that were s . There are also many conceivable ways to measure the quality of $p_{\text{approx}}(s)$ in approximating $p(s)$, but a common one and the one we choose to use here is the

KL-divergence, defined as follows:

$$\text{div}_{KL}(p||p_{\text{approx}}) \equiv \sum_{s \in \mathcal{S}} p(s) \log \frac{p(s)}{p_{\text{approx}}(s)}$$

By inspection, the KL-divergence is not defined unless $p_{\text{approx}}(s) = 0 \Rightarrow p(s) = 0$. It is clear that the KL-divergence is not defined for the empirical distribution until every possible sequence has been sampled at least once, which may take a very large number of samples. Therefore we use an alternate method for generating $p_{\text{approx}}(s)$ that is compatible with our choice of the KL-divergence for comparing $p_{\text{approx}}(s)$ to $p(s)$ — Multiplicative Weights. This algorithm is given below:

$MW_1(N, \eta) :$

Set $w_s \leftarrow 1 \forall s \in \mathcal{S}$

Repeat N times:

Sample $\sigma \leftarrow sample_1(\mathbf{u}, \varepsilon)$

Set $w_\sigma \leftarrow w_\sigma e^\eta$

Output $p_{\text{approx}}(s) \propto w_s$

In words, Multiplicative Weights initially assigns equal probability to all possible outputs, then samples from the mechanism N times, increasing the probabilities of outputs that occur at the expense of the probabilities of all other outputs. In order to do this, the update parameter η must be given, which controls how much the probability of each output changes when the output occurs. It is clear that MW_1 will always output an approximate probability distribution that leads to a defined KL-divergence for any true probability distribution.

2.2. Model 2

Let \mathcal{C} be the content universe. The set of all possible outputs is $\mathcal{S} \equiv \mathcal{C}^n$ — that is, an output is a sequence $s = (c_1, c_2, \dots, c_n)$ for $c_j \in \mathcal{C}$ and for some sequence length n . Define a graphical network $\mathcal{N} \equiv (\mathcal{C}, \mathcal{E})$ where \mathcal{E} is a multiset of directed edges with edge weights between nodes in \mathcal{C} . Specifically, an edge is of the form $\epsilon = (c_{\text{source}}, c_{\text{destination}}, \rho, u)$, where the source node c_{source} and the destination node $c_{\text{destination}}$ are both in \mathcal{C} , $\rho \in \mathcal{P}$ is a particular possible value for a preference that the user is allowed to choose at query-time defined by the content-creator, and $u \in (0, 1]$ is a utility defined by the content-creator. \mathcal{P} is the set of possible preferences at each node, which is assumed to be constant over the nodes for now — this can be interpreted as a set of options presented to the user when he arrives at a node. Here, \mathcal{E} can be considered the private database, as \mathcal{C} is assumed public knowledge. The content-creator is free to define as many or as few edges as she wishes.

The sampling mechanism is defined as follows:

$sample_2(\mathcal{N}, P, n, \varepsilon) :$

Set $c_1 \leftarrow c \in \mathcal{C}$ with probability $1/|\mathcal{C}|$

For $i = 2, \dots, n$:

Set $c_i \leftarrow c \in \mathcal{C}$ with probability $\propto \exp(\frac{1}{2}\varepsilon U(c|c_{i-1}))$

Output $s = (c_1, c_2, \dots, c_n)$

In this mechanism, $\mathcal{N} = (\mathcal{C}, \mathcal{E})$ is the content network, $P : \mathcal{C} \rightarrow \mathcal{P}$ is a mapping between content nodes and a user-given preference for each node, n is the sequence length, ε is the privacy parameter, and s is the output sequence. The utility function $U(c|c') = u \Leftrightarrow (c', c, P(c'), u) \in \mathcal{E}$, and $U(c|c') = 0 \Leftrightarrow (c', c, P(c'), u) \notin \mathcal{E} \forall u$. In words, if there is an edge defined in the content network that maps from node c' to node c and is labeled by the actual preference value the user gives for node c' , then the utility of this transition from c' to c is given by the utility of that edge. This corresponds to the transition being intended by the content-creator. Otherwise, the utility of the transition is zero. When we are at a node, the next node is picked based on the utilities of transitions away from that node by the exponential mechanism with privacy parameter ε (as before, the sensitivity of the utility function is 1).

The formulation of $sample_2$ given above has one input step followed by one output step. However, the mechanism can easily be reformulated into an online fashion by outputting each content node in the sequence sequentially, while probing for user preference at each step. This is the formulation one would use in an actual implementation.

Given the mapping of user inputs P , the probability of each sequence being output can be calculated exactly. We have $p(s|P) = p((c_1, c_2, \dots, c_n)|P) = p(c_1|P) \cdot p(c_2|c_1, P) \cdot p(c_3|c_2, P) \cdots p(c_n|c_{n-1}, P)$ because of the Markovianism of the mechanism. All of these conditional probabilities are known — they are simply the probabilities that appear in $sample_2$. Because $p(s|P)$ is known, for any $p_{\text{approx}}(s|P)$ we can use the KL-divergence again to compare them.

We suggest a similar Multiplicative Weights-based algorithm for the adversary to reconstruct $p_{\text{approx}}(s|P)$:

$MW_2(N, P, \eta) :$

Construct $\mathcal{N}_P \equiv (\mathcal{C}, \mathcal{E}_P)$ where

$\mathcal{E}_P = \{(c_{\text{source}}, c_{\text{dest}}, 1) \mid c_{\text{source}}, c_{\text{dest}} \in \mathcal{C}\}$

Repeat N times:

Sample $(c_1, \dots, c_n) \leftarrow sample_2(\mathcal{N}, P, n, \varepsilon)$

For $i = 1, \dots, n - 1$:

Set $(c_i, c_{i+1}, w) \leftarrow (c_i, c_{i+1}, w e^\eta)$ in \mathcal{E}_P

Let $p(c|c', P) = w_{c,c'} / \sum_\gamma w_{\gamma,c'}$

where $(c', d, w) \in \mathcal{E}_P \Leftrightarrow w_{d,c'} = w$

Output $p_{\text{approx}}((c_1, \dots, c_n)|P)$

$= \frac{1}{|\mathcal{C}|} p(c_2|c_1, P) \cdot p(c_3|c_2, P) \cdots p(c_n|c_{n-1}, P)$

This mechanism gives uniform probability to each transition between content nodes initially, but reweights probabilities based on the results of samples,

giving higher probabilities to transitions that it observes. The approximate probability of a sequence is just the product of the approximate transition probabilities that the adversary calculates (multiplied also by the probability of the first node being chosen as the starting node, which the adversary knows is uniform probability).

3. Results and Discussion

3.1. Model 1

We implemented Model 1 by picking a subset of the output space at random to be “intended” and assigned these random utilities in $(0, 1]$. We then investigated how the KL-divergence of the adversary’s approximate probability distribution over outputs relative to the true probability distribution over outputs change as a function of the number of queries to the mechanism as well as the value of the privacy parameter ε . The result is in Figure 1.




Fig. 1.—: The KL divergence of the adversary’s approximate probability distribution relative to the true probability distribution of Model 1, as a function of the number of queries to the mechanism as well as the privacy parameter ε . For the blue line, the probability a given sampled song was *intended* was 10%; for the green line it was 10%; for the red line it was 69%; for the cyan line it was 100%. The update parameter η was 10^{-2} .

Figure 1 demonstrates the opposite result from what we expected. Since a lower ε leads to more probability mass on unintended outputs — more noise — we expected the adversary to be hindered in his attempts to reconstruct the true probability distribution. Reality is quite the opposite — we see that the KL divergence starts out lower for lower ε . The reason for this is that the Multiplicative Weights algorithm starts by assigning equal probabilities to each output, which is the limit of the true probability dis-

tribution as $\varepsilon \rightarrow 0$. Therefore, for small ε a uniform probability distribution accurately approximates the true probability distribution.

A somewhat curious feature of the KL-divergence as a function of query count is that it does not monotonically-decrease. Instead, it seems to decrease, reach a minimum, then increase again. We were unable to explore more than 10^4 queries due to computing power limitations. This general feature of the KL-divergences was observed for various values of η . We would expect the approximate probability distribution to better approximate the true probability distribution as more queries were performed, so this is mysterious and warrants further investigation.

3.2. Model 2

We began by investigating the probabilities of various sequence outputs from *sample₂* for a model network — this is described in Figure 2. An interesting phenomenon is that plateaus can be seen in the sequence probability curves, indicating that many nodes have similar probabilities, but between these similar groups there is a sharp discontinuity in probability. This is most profoundly visible for lower fractions of defined links.




Fig. 2.— *Top*: Edge utilities vs. utility-ranking of edges in a content network with 10 nodes where each node-node edge has a chance of having nonzero utility given by the fraction of defined links. Utilities for edges with nonzero utilities were drawn from the uniform distribution over $(0, 1]$. This content network had $\mathcal{P} = \{0, 1\}$.

Bottom: Probabilities of various length-4 sequences in the above content network vs. the probability-ranking of those sequences. This probabilities were chosen by *sample₂* with an ε of 100.

We suspect, but have not verified, that the sequences on the same plateau have the same number of *intended* links in them — 3 for the first plateau, 2 for the second, 1 for the third, and 0 for the last. As the fraction of defined links increases, these plateaus are smoothed out as unintended nodes with utility 0 become intended nodes with random utility in $(0, 1]$.

A rough metric for quantifying the quality of the adversary’s reconstruction is by looking at the Top Sequences Error. The top $X\%$ error is defined as follows: take the top $X\%$ most probable sequences in reality and the top $X\%$ most probable sequences according to the adversary’s approximation. The proportion of true $X\%$ most probable sequences not found in the $X\%$ most probable sequences according to the adversary is the top $X\%$ sequences error. The Top Sequences Error is plotted as a function of the number of queries by an adversary using MW_2 with $\eta = 10^{-4}$ to reconstruct the network of Figure 2.




Fig. 3.—: Top Sequences Error vs. the number of queries by an adversary using MW_2 with $\eta = 10^{-4}$ to reconstruct the network of Figure 2.

For a reason we don’t understand, the top sequences error begins lower than would be expected from random chance, then rises — monotonically up to 10^7 queries for the top 10% error, but peaking and eventually declining for the top 50% error. We would expect that the top sequence error would decrease monotonically in the number of queries, but the actual behavior is much more complicated and warrants further investigation.

4. Conclusions

Ideally we would like to show that interactivity and noise combine to allow a linear increase in the effort of content-creators to lead to an exponential increase in the effort required to fully reconstruct a content-base. We were unable to show this as of yet. However, we have implemented a powerful model of content gener-

ation that is generalizable to many media and forms of content and reflects high-level principles of artistic media.

5. Further Work

There is much further work to be done in this area. A limitation of our Model 1 is that it does not support interactivity, even though it could be easily extended to do so. Further, we did not have a chance to explore in any detail the effect of interactivity in Model 2, but it is in a fully implemented state where, given more time, we would be able to show results. More methods for quantifying the ease of reconstruction of a content base should be explored, as well as alternative reconstruction strategies for the adversary. A “Model 3” might relax the Markovian property of Model 2 so that content-node transition probabilities can depend on the full sequence so far — this would provide a more general toolkit for content-production.

The biggest challenge the authors faced was defining exactly what could be protected from an adversary. In this paper, we settled with the attempting to protect the probabilities of various outputs, but we are not fully satisfied with this definition. More work should be put into exploring exactly what features of creative works can be protected from reconstruction without significant impacting the enjoyment of the work by authorized users. Ideally, a fully differentially-private mechanism could be designed to protect proprietary information so that the rich tools of the differential privacy literature could be employed.